

Взаимодействие между процессами

Process ID

```
#include <sys/types.h>
#include <unistd.h>
```

```
typedef int pid_t;
```

```
pid_t getpid();
pid_t getppid();
uid_t getuid();
gid_t getgid();
```

```
$ cat /proc/sys/kernel/pid_max
4194304
```

fork

[fork\(2\)](#)

```
#include <unistd.h>
```

```
pid_t fork();
```

Дочерний процесс получает копии:

- сегмента данных,
- кучи,
- стека.

Сегмент кода используется совместно.

[Dennis M. Ritchie, The Evolution of the Unix Time-sharing System](#)

```
#include <stdio.h>
#include <unistd.h>
```


```
int globvar = 6;
```

```
int main()
{
    int var;
    pid_t pid;

    var = 88;
    printf("перед вызовом функции fork\n");

    if ((pid = fork()) < 0) {
        perror("fork");
    } else if (pid == 0) {
        globvar++;
        var++;
    } else {
        sleep(2);
    }
    printf("pid = %ld, globvar = %d, var = %d\n",
        (long)getpid(), globvar, var);
    return 0;
}
```

[Copy-on-write](#)



Dirty CoW: [CVE-2016-5195](#)

Также наследуются:

- файловые дескрипторы (включая позицию),
- uid и gid,
- идентификаторы группы процессов и сеанса,
- управляющий терминал,
- cwd и root,
- umask,
- маска сигналов,
- переменные окружения,
- shared memory,
- отображения в память,
- ограничения на ресурсы.

Различаются:

- pid и ppid,
- значения `times(2)`,
- блокировки файлов (`fcntl(2)`) сбрасываются,
- таймеры сбрасываются,
- сигналы, ожидающие обработки, очищаются.

exec

`exec(3)`

```
#include <unistd.h>
int execl(const char* pathname, const char* arg0, ... /* (char*)0 */);
int execv(const char* pathname, char* const argv[]);
int execl(const char* pathname, const char* arg0, ...
    /* (char *)0, char* const envp[] */ );
int execve(const char* pathname, char* const argv[], char* const envp[]);
int execlp(const char* filename, const char* arg0, ... /* (char*)0 */ );
int execvp(const char* filename, char* const argv[]);
int fexecve(int fd, char* const argv[], char* const envp[]);
```

close-on-exec flag:

- `O_CLOEXEC` in `open(2)`

- FD_CLOEXEC in [fcntl\(2\)](#)

vfork

[vfork\(2\)](#)

The `vfork()` function has the same effect as `fork(2)`, except that the behavior is undefined if the process created by `vfork()` either modifies any data other than a variable of type `pid_t` used to store the return value from `vfork()`, or returns from the function in which `vfork()` was called, or calls any other function before successfully calling `_exit(2)` or one of the `exec(3)` family of functions.

CONFORMING TO: 4.3BSD; POSIX.1-2001 (but marked OBSOLETE). POSIX.1-2008 removes the specification of `vfork()`.

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int globvar = 6;

int main()
{
    int var;
    pid_t pid;

    var = 88;
    printf("перед вызовом функции vfork\n");
    if ((pid = vfork()) < 0) {
        perror("fork");
    } else if (pid == 0) {
        globvar++;
        var++;
        exit(0);
    }
    printf("pid = %ld, globvar = %d, var = %d\n",
        (long)getpid(), globvar, var);
    return 0;
}
```

[clone\(2\)](#), [rfork\(2\)](#)

[unshare\(2\)](#)

[posix_spawn\(3\)](#)

`fork()` is evil; `vfork()` is goodness; `fork()` would be better; `clone()` is stupid, перевод

system

[system\(3\)](#)

```
#include <stdlib.h>
```

```
int system(const char* cmdstring);
```

wait

[wait\(2\)](#)

```
#include <sys/wait.h>
```

```
pid_t wait(int* statloc);
pid_t waitpid(pid_t pid, int* statloc, int options);
int waitid(idtype_t idtype, id_t id, siginfo_t* infop, int options);
```

```
/* options */
```

```
WNOHANG;
WUNTRACED;
WCONTINUED;
```

```
/* status */
WIFEXITED;
WEXITSTATUS;
WIFSIGNALED;
WTERMSIG;
WCOUREDUMP;
WIFSTOPPED;
WSTOPSIG;
WIFCONTINUED;
```

```
/* idtype */
P_PID;
P_PGID;
P_ALL;
```

- Если родительский процесс завершается раньше дочернего, родителем последнего становится `init`.
- Если дочерний процесс завершается до того, как родитель вызывает `wait` или `waitpid` для его `pid`, такой процесс становится “зомби”.

```
#include <stdlib.h>
#include <sys/types.h>
#include <unistd.h>
```

```
int main ()
{
    pid_t child_pid;

    child_pid = fork ();
    if (child_pid > 0) {
        sleep (60);
    }
    else {
        exit (0);
    }
    return 0;
}
```

[Изучаем процессы в Linux](#)

Race condition

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
```

```
static void charatatime(char *str)
{
    char *ptr;
    int c;

    setbuf(stdout, NULL);
    for (ptr = str; (c = *ptr++) != 0; )
        putc(c, stdout);
}
```

```
int main()
{
    pid_t pid;
```

```

if ((pid = fork()) < 0) {
    perror("fork");
} else if (pid == 0) {
    charatotime("from child process\n");
} else {
    charatotime("from parent process\n");
}
exit(0);
}

```

Сигналы

signal(7)

- null signal: сигнал с номером 0.
- Сигналы, генерируемые терминалом (SIGINT, SIGHUP).
- Аппаратные ошибки (SIGSEGV, SIGBUS).
- Программные ситуации (SIGURG, SIGPIPE, SIGALRM).

```

/* ISO C99 signals. */
#define SIGINT      2      /* Interactive attention signal. */
#define SIGILL      4      /* Illegal instruction. */
#define SIGABRT     6      /* Abnormal termination. */
#define SIGFPE      8      /* Erroneous arithmetic operation. */
#define SIGSEGV     11     /* Invalid access to storage. */
#define SIGTERM     15     /* Termination request. */

/* Historical signals specified by POSIX. */
#define SIGHUP      1      /* Hangup. */
#define SIGQUIT     3      /* Quit. */
#define SIGTRAP     5      /* Trace/breakpoint trap. */
#define SIGKILL     9      /* Killed. */
#define SIGPIPE     13     /* Broken pipe. */
#define SIGALRM     14     /* Alarm clock. */

/* Historical signals specified by POSIX. */
#define SIGBUS      7      /* Bus error. */
#define SIGSYS      31     /* Bad system call. */

/* New(er) POSIX signals (1003.1-2008, 1003.1-2013). */
#define SIGURG      23     /* Urgent data is available at a socket. */
#define SIGSTOP     19     /* Stop, unblockable. */
#define SIGTSTP     20     /* Keyboard stop. */
#define SIGCONT     18     /* Continue. */
#define SIGCHLD     17     /* Child terminated or stopped. */
#define SIGTTIN     21     /* Background read from control terminal. */
#define SIGTTOU     22     /* Background write to control terminal. */
#define SIGPOLL     29     /* Pollable event occurred (System V). */
#define SIGXFSZ     25     /* File size limit exceeded. */
#define SIGXCPU     24     /* CPU time limit exceeded. */
#define SIGVTALRM   26     /* Virtual timer expired. */
#define SIGPROF     27     /* Profiling timer expired. */
#define SIGUSR1     10     /* User-defined signal 1. */
#define SIGUSR2     12     /* User-defined signal 2. */

```

```

#include <signal.h>

```

```

void psignal(int sig, const char* msg);
void psiginfo(const siginfo_t *pinfo, const char* msg);

```

```

kill(2), raise(3)

```

```
#include <signal.h>
```

```
int kill(pid_t pid, int sig);
```

```
int raise(int signo);
```

Default disposition

- Term
- Ign
- Core
- Stop
- Cont

```
signal(2)
```

```
#include <signal.h>
```

```
typedef void (*sighandler_t)(int);
```

```
sighandler_t signal(int signum, sighandler_t handler);
```

```
/* handler */
```

```
SIG_IGN;
```

```
SIG_DFL;
```

```
#include <stdio.h>
```

```
#include <signal.h>
```

```
#include <unistd.h>
```

```
static void sig_usr(int signo)
```

```
{
    if (signo == SIGUSR1)
        printf("принят сигнал SIGUSR1\n");
    else if (signo == SIGUSR2)
        printf("принят сигнал SIGUSR2\n");
    else
        printf("принят сигнал %d\n", signo);
}
```

```
int main()
```

```
{
    if (signal(SIGUSR1, sig_usr) == SIG_ERR)
        perror("signal SIGUSR1");
    if (signal(SIGUSR2, sig_usr) == SIG_ERR)
        perror("signal SIGUSR2");
    for(;;)
        pause();
}
```

```
pause(2)
```

```
#include <unistd.h>
```

```
int pause();
```

[Правила использования сигналов в Unix](#)

[О чем нельзя забывать при работе с POSIX-сигналами](#)

Прерванные сигналами системные вызовы

```
signal(7)
```

If a signal handler is invoked while a system call or library function call is blocked, then either: * the call is automatically restarted after the signal handler returns; or * the call fails with the error EINTR.

- read(2), write(2), ioctl(2) on "slow" device: terminal, pipe, socket.
- open(2) if it can block (e.g. for a FIFO).
- wait(2), waitpid(2).
- socket interfaces: accept(2), connect(2), send(2), recv(2).
- flock(2) and fcntl(2).
- POSIX message queue interfaces.
- getrandom(2).
- pthread_mutex_lock(3), pthread_cond_wait(3) and related.
- POSIX semaphore interfaces.
- read(2) from inotify(7).

again:

```
if ((n = read(fd, buf, BUFSIZE)) < 0) {
    if (errno == EINTR)
        goto again; /* просто прерванный системный вызов */
                    /* обработать другие возможные ошибки */
}
```

[The Rise of Worse is Better by Richard P. Gabriel, перевод](#)

Unix and C are the ultimate computer viruses.

Реентрабельные функции

signal-safety(7)

Fetching and setting the value of errno is async-signal-safe provided that the signal handler saves errno on entry and restores its value before returning.

alarm(2)

```
#include <unistd.h>
```

```
unsigned int alarm(unsigned int seconds);
```

getpwnam(3)

```
#include <stdio.h>
#include <string.h>
#include <signal.h>
#include <pwd.h>
#include <unistd.h>
```

```
static void my_alarm(int signo)
{
    struct passwd *rootptr;
    printf("внутри обработчика сигнала\n");
    if ((rootptr = getpwnam("root")) == NULL)
        perror("getpwnam");
    alarm(1);
}
```

```
static const char* username = "andrew";
```

```
int main(void)
{
    struct passwd* ptr;
    signal(SIGALRM, my_alarm);
    alarm(1);
    for (;;) {
        if ((ptr = getpwnam(username)) == NULL)
            perror("getpwnam");
        if (strcmp(ptr->pw_name, username) != 0)
```

```

        printf("возвращаемое значение повреждено! pw_name = %s\n",
               ptr->pw_name);
    }
}

```

Маска заблокированных сигналов

```

#include <signal.h>

int sigemptyset(sigset_t *set);
int sigfillset(sigset_t *set);
int sigaddset(sigset_t *set, int signo);
int sigdelset(sigset_t *set, int signo);

int sigismember(const sigset_t *set, int signo);
sigprocmask(2)
#include <signal.h>

int sigprocmask(int how, const sigset_t* set, sigset_t* oset);

/* how */
SIG_BLOCK;
SIG_UNBLOCK;
SIG_SETMASK;

```

sigaction

```

sigaction(2)
#include <signal.h>

int sigaction(int signo, const struct sigaction* act, struct sigaction* oact);

struct sigaction {
    void (*sa_handler)(int);
    void (*sa_sigaction)(int, siginfo_t *, void *);
    sigset_t sa_mask;
    int sa_flags;
};

```

Межпроцессовые коммуникации

Межпроцессовые коммуникации LINUX

Полудуплексные каналы (pipe)

```

ls | sort

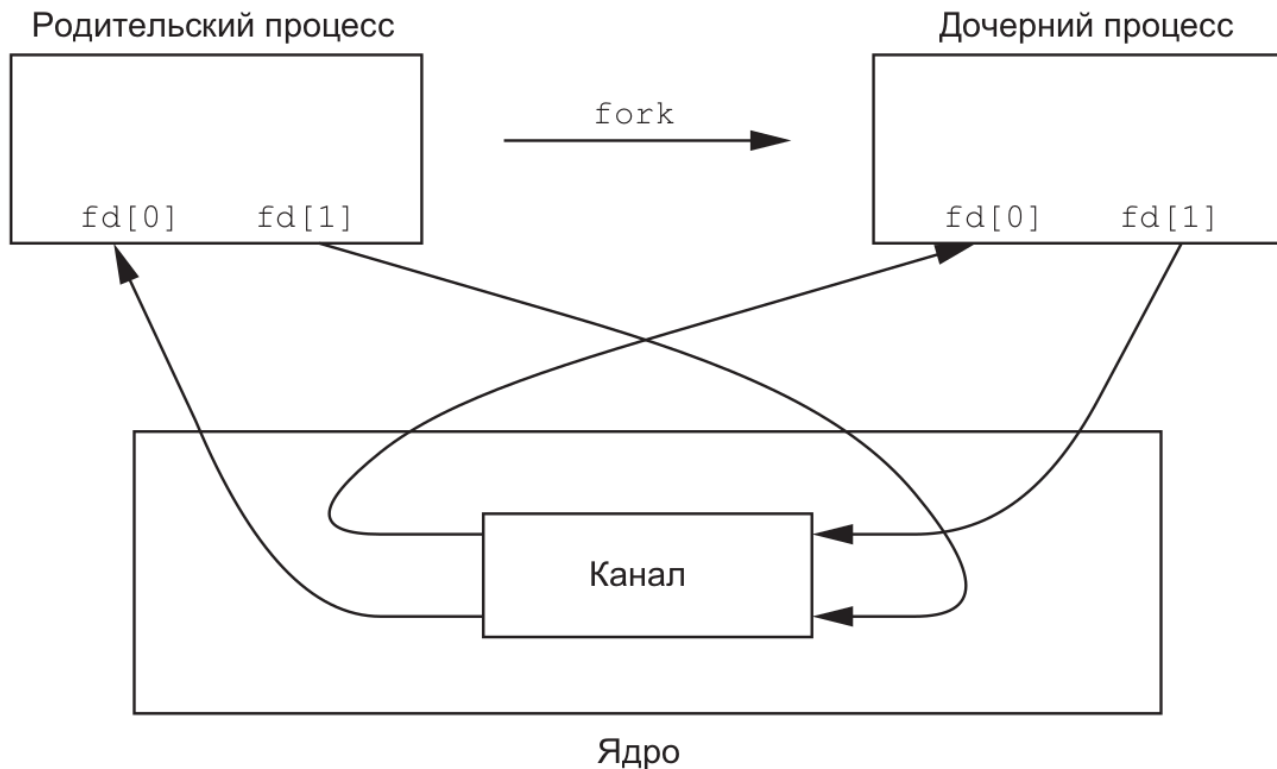
pipe(2)
#include <unistd.h>

int
pipe(int pipefd[2]);

int
pipe2(int pipefd[2], int flags);

#define PIPE_BUF 4096 /* Linux */

```

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>

int main(void)
{
    int    fd[2], nbytes;
    pid_t  childpid;
    char    string[] = "Hello, world!\n";
    char    readbuffer[80];

    pipe(fd);
    if ((childpid = fork()) == -1)
    {
        perror("fork");
        exit(1);
    }
    if (childpid == 0)
    {
        /* Child process closes up input side of pipe */
        close(fd[0]);
        /* Send "string" through the output side of pipe */
        write(fd[1], string, strlen(string));
        exit(0);
    }
    else
    {
        /* Parent process closes up output side of pipe */
        close(fd[1]);
        /* Read in a string from the pipe */
        nbytes = read(fd[0], readbuffer, sizeof(readbuffer));
        printf("Received string: %s", readbuffer);
    }
    return(0);
}
```

pipe(7)

A pipe has a limited capacity. If the pipe is full, then a write(2) will block or fail, depending on whether the O_NONBLOCK flag is set. Different implementations have different limits for the pipe capacity. Applications should not rely on a particular capacity: an application should be designed so that a reading process consumes data as soon as it is available, so that a writing process does not remain blocked.

Since Linux 2.6.11, the pipe capacity is 16 pages (i.e., 65,536 bytes in a system with a page size of 4096 bytes).

Так все же, насколько быстры каналы в Linux?

Benchmarks for Inter-Process-Communication Techniques

popen

popen(3)

```
#include <stdio.h>
```

```
FILE *popen(const char* cmdstring, const char* type);
```

```
int pclose(FILE *stream);
```

Именованные каналы (FIFO)

```
$ mknod MYFIFO p
```

mknod(2)

```
#include <sys/stat.h>
```

```
#include <fcntl.h>
```

```
#include <unistd.h>
```

```
int mknod(const char *pathname, mode_t mode, dev_t dev);
```

```
#include <stdio.h>
```

```
#include <stdlib.h>
```

```
#include <sys/stat.h>
```

```
#include <unistd.h>
```

```
#include <linux/stat.h>
```

```
#define FIFO_FILE      "MYFIFO"
```

```
int main(void)
```

```
{
```

```
    FILE *fp;
```

```
    char readbuf[80];
```

```
    /* Create the FIFO if it does not exist */
```

```
    umask(0);
```

```
    mknod(FIFO_FILE, S_IFIFO|0666, 0);
```

```
    while(1)
```

```
    {
```

```
        fp = fopen(FIFO_FILE, "r");
```

```
        fgets(readbuf, 80, fp);
```

```
        printf("Received string: %s\n", readbuf);
```

```
        fclose(fp);
```

```
    }
```

```
    return(0);
```

```
}
```

UNIX domain sockets

unix(7)

unix domain sockets vs. internet sockets

```
#include <sys/socket.h>
```

```
int socketpair(int domain, int type, int protocol, int sockfd[2]);
```

```
/* Linux, Solaris */
```

```
struct sockaddr_un {  
    sa_family_t sun_family; /* AF_UNIX */  
    char sun_path[108];     /* полное имя */  
};
```

```
/* FreeBSD, MacOS */
```

```
struct sockaddr_un {  
    unsigned char sun_len; /* длина адреса сокета */  
    sa_family_t sun_family; /* AF_UNIX */  
    char sun_path[104];     /* полное имя */  
};
```

```
#include <stdio.h>
```

```
#include <stddef.h>
```

```
#include <sys/socket.h>
```

```
#include <sys/un.h>
```

```
int main()
```

```
{  
    int fd, size;  
    struct sockaddr_un un;  
  
    un.sun_family = AF_UNIX;  
    strcpy(un.sun_path, "foo.socket");  
    if ((fd = socket(AF_UNIX, SOCK_STREAM, 0)) < 0)  
        perror("socket");  
    size = offsetof(struct sockaddr_un, sun_path) + strlen(un.sun_path);  
    if (bind(fd, (struct sockaddr *)&un, size) < 0)  
        perror("bind");  
  
    printf("Done!\n");  
    return 0;  
}
```

UDS server

```
#include <stdio.h>
```

```
#include <stdlib.h>
```

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
#include <sys/un.h>
```

```
#include <unistd.h>
```

```
#define NAME "echo.socket"
```

```
int main()
```

```
{  
    int sock, msgsock, rval;  
    struct sockaddr_un server;  
    char buf[1024];  
  
    sock = socket(AF_UNIX, SOCK_STREAM, 0);
```

```

if (sock < 0) {
    perror("opening stream socket");
    exit(1);
}
server.sun_family = AF_UNIX;
strcpy(server.sun_path, NAME);
if (bind(sock, (struct sockaddr *) &server, sizeof(struct sockaddr_un))) {
    perror("binding stream socket");
    exit(1);
}
printf("Socket has name %s\n", server.sun_path);
listen(sock, 5);
for (;;) {
    msgsock = accept(sock, 0, 0);
    if (msgsock == -1)
    {
        perror("accept");
        break;
    }
    else do {
        bzero(buf, sizeof(buf));
        if ((rval = read(msgsock, buf, 1024)) < 0)
            perror("reading stream message");
        else if (rval == 0)
            printf("Ending connection\n");
        else
            send(msgsock, buf, strlen(buf), 0);
    } while (rval > 0);
    close(msgsock);
}
close(sock);
unlink(NAME);
}

```

[Windows/WSL Interop with AF_UNIX](#)